

## ✓ Pytorch with the MNIST Dataset - MINST

```
## import libraries
from __future__ import print_function
import argparse
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision

import torch.optim as optim
from torchvision import datasets, transforms
from torch.autograd import Variable

print(torch.__version__)

↩ 2.5.0+cu121

args={}
kwargs={}
args['batch_size']=32
args['test_batch_size']=32
args['epochs']=1 #The number of Epochs is the number of times you go through the full dataset.
args['lr']=0.01 #Learning rate is how fast it will descend.
args['momentum']=0.5 #SGD momentum (default: 0.5) Momentum is a moving average of our gradients (helps to keep direction).

args['seed']=1 #random seed
args['log_interval']=10
args['cuda']=True #if the computer has a GPU, type True, otherwise, False
```

This code is adopted from the pytorch examples repository. It is licensed under BSD 3-Clause "New" or "Revised" License. Source: <https://github.com/pytorch/examples/> LICENSE: <https://github.com/pytorch/examples/blob/master/LICENSE>

## ✓ Load Dataset

The first step before training the model is to import the data. We will use the [MNIST dataset](#) which is like the Hello World dataset of machine learning.

Besides importing the data, we will also do a few more things:

- We will transform the data into tensors using the `transforms` module
- We will use `DataLoader` to build convenient data loaders or what are referred to as iterators, which makes it easy to efficiently feed data in batches to deep learning models.
- As hinted above, we will also create batches of the data by setting the `batch` parameter inside the data loader. Notice we use batches of 32 in this tutorial but you can change it to 64 if you like. I encourage you to experiment with different batches.

```
## transformations
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

## download and load training dataset
trainset = datasets.MNIST(root='../data', train=True, download=True, transform=transform)
train_loader = torch.utils.data.DataLoader(trainset, batch_size=args['batch_size'], shuffle=True, **kwargs)

## download and load testing dataset
testset = torchvision.datasets.MNIST(root='../data', train=False, download=True, transform=transform)
test_loader = torch.utils.data.DataLoader(testset, batch_size=args['test_batch_size'], shuffle=True, **kwargs)
```



```

<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz to ../data/MNIST/raw/t10k-labels-idx1-ubyte.
100%|██████████| 4.54k/4.54k [00:00<00:00, 2.83MB/s]
Extracting ../data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ../data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Failed to download (trying next):
<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz to ../data/MNIST/raw/train-images-idx3-ubyte.
100%|██████████| 9.91M/9.91M [00:00<00:00, 16.3MB/s]
Extracting ../data/MNIST/raw/train-images-idx3-ubyte.gz to ../data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Failed to download (trying next):
<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz to ../data/MNIST/raw/train-labels-idx1-ubyte.
100%|██████████| 28.9k/28.9k [00:00<00:00, 501kB/s]
Extracting ../data/MNIST/raw/train-labels-idx1-ubyte.gz to ../data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Failed to download (trying next):
<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz to ../data/MNIST/raw/t10k-images-idx3-ubyte.g
100%|██████████| 1.65M/1.65M [00:00<00:00, 4.54MB/s]
Extracting ../data/MNIST/raw/t10k-images-idx3-ubyte.gz to ../data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Failed to download (trying next):
<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz to ../data/MNIST/raw/t10k-labels-idx1-ubyte.g
100%|██████████| 4.54k/4.54k [00:00<00:00, 6.83MB/s]Extracting ../data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ../data/MNIST/raw

```

## ✓ Exploring the Data

As a practitioner and researcher, I am always spending a bit of time and effort exploring and understanding the dataset. It's fun and this is a good practise to ensure that everything is in order.

Let's check what the train and test dataset contains. I will use `matplotlib` to print out some of the images from our dataset.

```

import matplotlib.pyplot as plt
import numpy as np


## functions to show an image
def imshow(img):
    #img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))

## get some random training images
dataiter = iter(train_loader)
images, labels = next(dataiter)

## show images
imshow(torchvision.utils.make_grid(images))

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)



Let's check the dimensions of a batch.

```
for images, labels in train_loader:
    print("Image batch dimensions:", images.shape)
    print("Image label dimensions:", labels.shape)
    break
```

Image batch dimensions: torch.Size([32, 1, 28, 28])  
Image label dimensions: torch.Size([32])

## ✓ The Model

We provide two fully-connected neural net as the initial architecture.

Here are a few notes for those who are beginning with PyTorch:

- The model below consists of an `__init__()` portion which is where you include the layers and components of the neural network. In our model, we have two fully-connected network. We are dealing with an image dataset that is in a grayscale so we only need one channel going in, hence `in_channels=1`.
- After the first layer, we also apply an activation function such as `ReLU`. For prediction purposes, we then apply a `softmax` layer to the last transformation and return the output of that.

```
class Net(nn.Module):
    #This defines the structure of the NN.
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 25, 3) #1x28x28 --> 25x26x26
        self.pool = nn.MaxPool2d(2, 2) #25x26x26 --> 25x13x13
        self.fc1 = nn.Linear(25 * 13 * 13, 256)
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(x)
        x=x.view(-1,25 * 13 * 13)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        #Softmax gets probabilities.
        return F.log_softmax(x, dim=1)
```

**Now, add one CNN layer with a pooling to the above neural network and rerun the code to see whether you get higher prediction accuracy on the test set.**

For example, you may try `out_channels=32`. Kernel size is 5, and for the rest of parameters we use the default values which you can find [here](#).

- In short, the convolutional layer transforms the input data into a specific dimension that has to be considered in the linear layer.

Make sure your flatten the output of CNN layer excluding # of batch so that the input of each example/batch has the same size of the first neural net.

Tips: You can use `x.view(-1, # of input size of the first fully-connected layer)` or you can use `torch.flatten(x, 1)`.

I always encourage to test the model with 1 batch to ensure that the output dimensions are what we expect.

Start coding or [generate](#) with AI.

```
## test the model with 1 batch
model = Net()
#print(model)
for images, labels in train_loader:
    print("batch size:", args['batch_size'])
    out = model(images)
    print(out.shape)
    break
```

```
batch size: 32
torch.Size([32, 10])
```

## ✓ Training the Model

Now we are ready to train the model.

```
def train(epoch):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        if args['cuda']:
            data, target = data.cuda(), target.cuda()
            #Variables in Pytorch are differentiable.
            data, target = Variable(data), Variable(target)
            #This will zero out the gradients for this batch.
            optimizer.zero_grad()
            output = model(data)
            # Calculate the loss The negative log likelihood loss. It is useful to train a classification problem with C classes.
            loss = F.nll_loss(output, target)
            #dloss/dx for every Variable
            loss.backward()
            #to do a one-step update on our parameter.
            optimizer.step()
            #Print out the loss periodically.
            if batch_idx % args['log_interval'] == 0:
                print('Train Epoch: {} [{}/{}] ({:.0f}%) \tLoss: {:.6f}'.format(
                    epoch, batch_idx * len(data), len(train_loader.dataset),
                    100. * batch_idx / len(train_loader), loss.data.item()))
```

```
def test():
    model.eval()
    test_loss = 0
    correct = 0

    with torch.no_grad():
        for data, target in test_loader:
            if args['cuda']:
                data, target = data.cuda(), target.cuda()
                data, target = Variable(data), Variable(target)
                output = model(data)
                test_loss += F.nll_loss(output, target, size_average=False).data.item() # sum up batch loss
                pred = output.data.max(1, keepdim=True)[1] # get the index of the max log-probability
                correct += pred.eq(target.data.view_as(pred)).long().cpu().sum()

    test_loss /= len(test_loader.dataset)
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%) \n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))
```

```
model = Net()
if args['cuda']:
    model.cuda()

optimizer = optim.SGD(model.parameters(), lr=args['lr'], momentum=args['momentum'])

for epoch in range(1, args['epochs'] + 1):
    train(epoch)
    test()
```

```
Train Epoch: 1 [0/60000 (0%)] Loss: 2.372720
Train Epoch: 1 [320/60000 (1%)] Loss: 1.802728
Train Epoch: 1 [640/60000 (1%)] Loss: 1.231960
Train Epoch: 1 [960/60000 (2%)] Loss: 0.903316
Train Epoch: 1 [1280/60000 (2%)] Loss: 0.557140
Train Epoch: 1 [1600/60000 (3%)] Loss: 0.635397
Train Epoch: 1 [1920/60000 (3%)] Loss: 0.507572
Train Epoch: 1 [2240/60000 (4%)] Loss: 0.887186
Train Epoch: 1 [2560/60000 (4%)] Loss: 0.572372
Train Epoch: 1 [2880/60000 (5%)] Loss: 0.420833
Train Epoch: 1 [3200/60000 (5%)] Loss: 0.218439
```

Train Epoch: 1	[3520/60000 (6%)]	Loss: 0.237025
Train Epoch: 1	[3840/60000 (6%)]	Loss: 0.367036
Train Epoch: 1	[4160/60000 (7%)]	Loss: 0.186796
Train Epoch: 1	[4480/60000 (7%)]	Loss: 0.315542
Train Epoch: 1	[4800/60000 (8%)]	Loss: 0.234172
Train Epoch: 1	[5120/60000 (9%)]	Loss: 0.526702
Train Epoch: 1	[5440/60000 (9%)]	Loss: 0.429692
Train Epoch: 1	[5760/60000 (10%)]	Loss: 0.461097
Train Epoch: 1	[6080/60000 (10%)]	Loss: 0.884914
Train Epoch: 1	[6400/60000 (11%)]	Loss: 0.434885
Train Epoch: 1	[6720/60000 (11%)]	Loss: 0.315350
Train Epoch: 1	[7040/60000 (12%)]	Loss: 0.360749
Train Epoch: 1	[7360/60000 (12%)]	Loss: 0.348866
Train Epoch: 1	[7680/60000 (13%)]	Loss: 0.361742
Train Epoch: 1	[8000/60000 (13%)]	Loss: 0.388472
Train Epoch: 1	[8320/60000 (14%)]	Loss: 0.422414
Train Epoch: 1	[8640/60000 (14%)]	Loss: 0.493869
Train Epoch: 1	[8960/60000 (15%)]	Loss: 0.304426
Train Epoch: 1	[9280/60000 (15%)]	Loss: 0.163742
Train Epoch: 1	[9600/60000 (16%)]	Loss: 0.192996
Train Epoch: 1	[9920/60000 (17%)]	Loss: 0.327874
Train Epoch: 1	[10240/60000 (17%)]	Loss: 0.125876
Train Epoch: 1	[10560/60000 (18%)]	Loss: 0.406872
Train Epoch: 1	[10880/60000 (18%)]	Loss: 0.193155
Train Epoch: 1	[11200/60000 (19%)]	Loss: 0.404689
Train Epoch: 1	[11520/60000 (19%)]	Loss: 0.561336
Train Epoch: 1	[11840/60000 (20%)]	Loss: 0.099842
Train Epoch: 1	[12160/60000 (20%)]	Loss: 0.095993
Train Epoch: 1	[12480/60000 (21%)]	Loss: 0.374357
Train Epoch: 1	[12800/60000 (21%)]	Loss: 0.368063
Train Epoch: 1	[13120/60000 (22%)]	Loss: 0.233614
Train Epoch: 1	[13440/60000 (22%)]	Loss: 0.490324
Train Epoch: 1	[13760/60000 (23%)]	Loss: 0.376398
Train Epoch: 1	[14080/60000 (23%)]	Loss: 0.282401
Train Epoch: 1	[14400/60000 (24%)]	Loss: 0.211902
Train Epoch: 1	[14720/60000 (25%)]	Loss: 0.257408
Train Epoch: 1	[15040/60000 (25%)]	Loss: 0.365946
Train Epoch: 1	[15360/60000 (26%)]	Loss: 0.186456
Train Epoch: 1	[15680/60000 (26%)]	Loss: 0.370927
Train Epoch: 1	[16000/60000 (27%)]	Loss: 0.215926
Train Epoch: 1	[16320/60000 (27%)]	Loss: 0.097077
Train Epoch: 1	[16640/60000 (28%)]	Loss: 0.232138
Train Epoch: 1	[16960/60000 (28%)]	Loss: 0.696414
Train Epoch: 1	[17280/60000 (29%)]	Loss: 0.135584
Train Epoch: 1	[17600/60000 (29%)]	Loss: 0.324681
Train Epoch: 1	[17920/60000 (30%)]	Loss: 0.101354

Start coding or [generate](#) with AI.