

✓ Pytorch with the MNIST Dataset - MINST

```
## import libraries
from __future__ import print_function
import argparse
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision

import torch.optim as optim
from torchvision import datasets, transforms
from torch.autograd import Variable

print(torch.__version__)

2.5.0+cu121

args={}
kwargs={}
args['batch_size']=32
args['test_batch_size']=32
args['epochs']=1 #The number of Epochs is the number of times you go through the full dataset.
args['lr']=0.01 #Learning rate is how fast it will descend.
args['momentum']=0.5 #SGD momentum (default: 0.5) Momentum is a moving average of our gradients (helps to keep direction).

args['seed']=1 #random seed
args['log_interval']=10
args['cuda']=True #if the computer has a GPU, type True, otherwise, False
```

This code is adopted from the pytorch examples repository. It is licensed under BSD 3-Clause "New" or "Revised" License. Source: <https://github.com/pytorch/examples/> LICENSE: <https://github.com/pytorch/examples/blob/master/LICENSE>

✓ Load Dataset

The first step before training the model is to import the data. We will use the [MNIST dataset](#) which is like the Hello World dataset of machine learning.

Besides importing the data, we will also do a few more things:

- We will transform the data into tensors using the `transforms` module
- We will use `DataLoader` to build convenient data loaders or what are referred to as iterators, which makes it easy to efficiently feed data in batches to deep learning models.
- As hinted above, we will also create batches of the data by setting the `batch` parameter inside the data loader. Notice we use batches of 32 in this tutorial but you can change it to 64 if you like. I encourage you to experiment with different batches.

```
## transformations
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

## download and load training dataset
trainset = datasets.MNIST(root='../data', train=True, download=True, transform=transform)
train_loader = torch.utils.data.DataLoader(trainset, batch_size=args['batch_size'], shuffle=True, **kwargs)

## download and load testing dataset
testset = torchvision.datasets.MNIST(root='../data', train=False, download=True, transform=transform)
test_loader = torch.utils.data.DataLoader(testset, batch_size=args['test_batch_size'], shuffle=True, **kwargs)
```

✓ Exploring the Data

As a practitioner and researcher, I am always spending a bit of time and effort exploring and understanding the dataset. It's fun and this is a good practise to ensure that everything is in order.

Let's check what the train and test dataset contains. I will use `matplotlib` to print out some of the images from our dataset.

```
import matplotlib.pyplot as plt
```


```
import numpy as np
```

```
## functions to show an image
def imshow(img):
    #img = img / 2 + 0.5    # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
```

```
## get some random training images
dataiter = iter(train_loader)
images, labels = next(dataiter)
```

```
## show images
imshow(torchvision.utils.make_grid(images))
```

⚠ WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)



Let's check the dimensions of a batch.

```
for images, labels in train_loader:
    print("Image batch dimensions:", images.shape)
    print("Image label dimensions:", labels.shape)
    break
```

⚠ Image batch dimensions: torch.Size([64, 1, 28, 28])
Image label dimensions: torch.Size([64])

✓ The Model

We provide two fully-connected neural net as the initial architecture.

Here are a few notes for those who are beginning with PyTorch:

- The model below consists of an `__init__()` portion which is where you include the layers and components of the neural network. In our model, we have two fully-connected network. We are dealing with an image dataset that is in a grayscale so we only need one channel going in, hence `in_channels=1`.
- After the first layer, we also apply an activation function such as `ReLU`. For prediction purposes, we then apply a `softmax` layer to the last transformation and return the output of that.

```
class Net(nn.Module):
    #This defines the structure of the NN.
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 16, 3) #1x28x28 --> 16x26x26
        self.pool = nn.MaxPool2d(2, 2) #16x26x26 --> 16x13x13
        self.fc1 = nn.Linear(16 * 13 * 13, 256)
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(x)
        x=x.view(-1,16 * 13 * 13)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        #Softmax gets probabilities.
        return F.log_softmax(x, dim=1)
```

Now, add one CNN layer with a pooling to the above neural network and rerun the code to see whether you get higher prediction accuracy on the test set.

For example, you may try `out_channels=32`. Kernel size is 5, and for the rest of parameters we use the default values which you can find [here](#).

- In short, the convolutional layer transforms the input data into a specific dimension that has to be considered in the linear layer.

Make sure you flatten the output of CNN layer excluding # of batch so that the input of each example/batch has the same size of the first neural net.

Tips: You can use `x.view(-1, # of input size of the first fully-connected layer)` or you can use `torch.flatten(x, 1)`.

I always encourage to test the model with 1 batch to ensure that the output dimensions are what we expect.

Start coding or [generate](#) with AI.

```
## test the model with 1 batch
model = Net()
#print(model)
for images, labels in train_loader:
    print("batch size:", args['batch_size'])
    out = model(images)
    print(out.shape)
    break
```

```
batch size: 32
torch.Size([64, 10])
```

✓ Training the Model

Now we are ready to train the model.

```
def train(epoch):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        if args['cuda']:
            data, target = data.cuda(), target.cuda()
        #Variables in Pytorch are differentiable.
        data, target = Variable(data), Variable(target)
        #This will zero out the gradients for this batch.
        optimizer.zero_grad()
        output = model(data)
        # Calculate the loss The negative log likelihood loss. It is useful to train a classification problem with C classes.
        loss = F.nll_loss(output, target)
        #dloss/dx for every Variable
        loss.backward()
        #to do a one-step update on our parameter.
        optimizer.step()
        #Print out the loss periodically.
        if batch_idx % args['log_interval'] == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.data.item()))
```

```
def test():
    model.eval()
    test_loss = 0
    correct = 0

    with torch.no_grad():
        for data, target in test_loader:
            if args['cuda']:
                data, target = data.cuda(), target.cuda()
            data, target = Variable(data), Variable(target)
            output = model(data)
            test_loss += F.nll_loss(output, target, size_average=False).data.item() # sum up batch loss
            pred = output.data.max(1, keepdim=True)[1] # get the index of the max log-probability
            correct += pred.eq(target.data.view_as(pred)).long().cpu().sum()

    test_loss /= len(test_loader.dataset)
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)\n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))
```

```

model = Net()
if args['cuda']:
    model.cuda()

optimizer = optim.SGD(model.parameters(), lr=args['lr'], momentum=args['momentum'])

for epoch in range(1, args['epochs'] + 1):
    train(epoch)
    test()

```

Train Epoch: 1	Loss: 0.220000
Train Epoch: 1 [27520/60000 (46%)]	Loss: 0.209637
Train Epoch: 1 [28160/60000 (47%)]	Loss: 0.339335
Train Epoch: 1 [28800/60000 (48%)]	Loss: 0.347948
Train Epoch: 1 [29440/60000 (49%)]	Loss: 0.298498
Train Epoch: 1 [30080/60000 (50%)]	Loss: 0.271223
Train Epoch: 1 [30720/60000 (51%)]	Loss: 0.216389
Train Epoch: 1 [31360/60000 (52%)]	Loss: 0.199735
Train Epoch: 1 [32000/60000 (53%)]	Loss: 0.283339
Train Epoch: 1 [32640/60000 (54%)]	Loss: 0.088744
Train Epoch: 1 [33280/60000 (55%)]	Loss: 0.266086
Train Epoch: 1 [33920/60000 (57%)]	Loss: 0.176704
Train Epoch: 1 [34560/60000 (58%)]	Loss: 0.392657
Train Epoch: 1 [35200/60000 (59%)]	Loss: 0.111579
Train Epoch: 1 [35840/60000 (60%)]	Loss: 0.139564
Train Epoch: 1 [36480/60000 (61%)]	Loss: 0.206128
Train Epoch: 1 [37120/60000 (62%)]	Loss: 0.157328
Train Epoch: 1 [37760/60000 (63%)]	Loss: 0.351994
Train Epoch: 1 [38400/60000 (64%)]	Loss: 0.415341
Train Epoch: 1 [39040/60000 (65%)]	Loss: 0.338110
Train Epoch: 1 [39680/60000 (66%)]	Loss: 0.407146
Train Epoch: 1 [40320/60000 (67%)]	Loss: 0.425015
Train Epoch: 1 [40960/60000 (68%)]	Loss: 0.305363
Train Epoch: 1 [41600/60000 (69%)]	Loss: 0.472191
Train Epoch: 1 [42240/60000 (70%)]	Loss: 0.340229
Train Epoch: 1 [42880/60000 (71%)]	Loss: 0.244672
Train Epoch: 1 [43520/60000 (72%)]	Loss: 0.160745
Train Epoch: 1 [44160/60000 (74%)]	Loss: 0.268718
Train Epoch: 1 [44800/60000 (75%)]	Loss: 0.304855
Train Epoch: 1 [45440/60000 (76%)]	Loss: 0.241072
Train Epoch: 1 [46080/60000 (77%)]	Loss: 0.257948
Train Epoch: 1 [46720/60000 (78%)]	Loss: 0.194221
Train Epoch: 1 [47360/60000 (79%)]	Loss: 0.244654
Train Epoch: 1 [48000/60000 (80%)]	Loss: 0.180502
Train Epoch: 1 [48640/60000 (81%)]	Loss: 0.230826
Train Epoch: 1 [49280/60000 (82%)]	Loss: 0.172106
Train Epoch: 1 [49920/60000 (83%)]	Loss: 0.409677
Train Epoch: 1 [50560/60000 (84%)]	Loss: 0.255971
Train Epoch: 1 [51200/60000 (85%)]	Loss: 0.178604
Train Epoch: 1 [51840/60000 (86%)]	Loss: 0.258945
Train Epoch: 1 [52480/60000 (87%)]	Loss: 0.195263
Train Epoch: 1 [53120/60000 (88%)]	Loss: 0.282956
Train Epoch: 1 [53760/60000 (90%)]	Loss: 0.121762
Train Epoch: 1 [54400/60000 (91%)]	Loss: 0.190646
Train Epoch: 1 [55040/60000 (92%)]	Loss: 0.159668
Train Epoch: 1 [55680/60000 (93%)]	Loss: 0.183767
Train Epoch: 1 [56320/60000 (94%)]	Loss: 0.070564
Train Epoch: 1 [56960/60000 (95%)]	Loss: 0.295069
Train Epoch: 1 [57600/60000 (96%)]	Loss: 0.303115
Train Epoch: 1 [58240/60000 (97%)]	Loss: 0.196017
Train Epoch: 1 [58880/60000 (98%)]	Loss: 0.198433
Train Epoch: 1 [59520/60000 (99%)]	Loss: 0.110093

/usr/local/lib/python3.10/dist-packages/torch/nn/_reduction.py:51: UserWarning: size_average and reduce args will be deprecated, warnings.warn(warning.format(ret))

Test set: Average loss: 0.1699, Accuracy: 9485/10000 (95%)

Start coding or [generate](#) with AI.

